

Going deeper with dplyr

王小二 20190001

2019-10-04

1 Loading dplyr and the nycflights13 dataset

```
# load packages
suppressMessages(library(dplyr))
library(nycflights13)

# print the flights dataset from nycflights13
flights
```



```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2 Choosing columns: select, rename

```
# besides just using select() to pick columns...
```

```
flights %>% select(carrier, flight)
```

```
## # A tibble: 336,776 x 2
```

```
##   carrier flight
```

```
##   <chr>    <int>
```

```
## 1 UA      1545
```

```
## 2 UA      1714
```

```
## 3 AA      1141
```

```
## 4 B6       725
```

```
## 5 DL       461
```

```
## 6 UA      1696
```

```
## 7 B6       507
```

```
## 8 EV      5708
```

```
## 9 B6        79
```

```
## 10 AA      301
```

```
## # ... with 336,766 more rows
```

```
# ...you can use the minus sign to hide columns
```

```
flights %>% select(-month, -day)
```

```
## # A tibble: 336,776 x 17
```

```
##   year dep_time sched_dep_time dep_delay arr_time sched_arr_time
```

```
##   <int>   <int>         <int>    <dbl>   <int>         <int>
```

```
## 1 2013     517           515        2     830           819
```

```
## 2 2013     533           529        4     850           830
```

```
## 3 2013     542           540        2     923           850
```

```
## 4 2013     544           545       -1    1004          1022
```

```
## 5 2013     554           600       -6     812           837
```

```
## 6 2013     554           558       -4     740           728
```

```
## 7 2013     555           600       -5     913           854
```

```
## 8 2013     557           600       -3     709           723
```

```
## 9 2013     557           600       -3     838           846
```

```
## 10 2013     558           600       -2     753           745
```

```
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

```
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
```

```
## #   time_hour <dtm>
```

```
# hide a range of columns
```

```
flights %>% select(-(dep_time:arr_delay))
```

```
# hide any column with a matching name
```

```
flights %>% select(-contains("time"))
```

```
# pick columns using a character vector of column names
```

```
cols <- c("carrier", "flight", "tailnum")
```

```
flights %>% select(one_of(cols))
```

```
## # A tibble: 336,776 x 3
```

```
##   carrier flight tailnum
```

```
##   <chr>      <int> <chr>
```

```
## 1 UA         1545 N14228
```

```
## 2 UA         1714 N24211
```

```
## 3 AA         1141 N619AA
```

```
## 4 B6          725 N804JB
```

```
## 5 DL          461 N668DN
```

```
## 6 UA         1696 N39463
```

```
## 7 B6          507 N516JB
```

```
## 8 EV         5708 N829AS
```

```
## 9 B6           79 N593JB
```

```
## 10 AA         301 N3ALAA
```

```
## # ... with 336,766 more rows
```

```
# select() can be used to rename columns, though all columns not mentioned are dropped
```

```
flights %>% select(tail = tailnum)
```

```
## # A tibble: 336,776 x 1
```

```
##   tail
```

```
##   <chr>
```

```
## 1 N14228
```

```
## 2 N24211
```

```
## 3 N619AA
```

```
## 4 N804JB
```

```
## 5 N668DN
```

```
## 6 N39463
```

```
## 7 N516JB
```

```
## 8 N829AS
```

```
## 9 N593JB
## 10 N3ALAA
## # ... with 336,766 more rows
```

```
# rename() does the same thing, except all columns not mentioned are kept
flights %>% rename(tail = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tail <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3 Choosing rows: filter, between, slice, sample_n, top_n, distinct

```
# filter() supports the use of multiple conditions
flights %>% filter(dep_time >= 600, dep_time <= 605)
```

```
## # A tibble: 2,460 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     600           600         0     851
## 2  2013     1     1     600           600         0     837
## 3  2013     1     1     601           600         1     844
## 4  2013     1     1     602           610        -8     812
```

```
## 5 2013 1 1 602 605 -3 821
## 6 2013 1 2 600 600 0 814
## 7 2013 1 2 600 605 -5 751
## 8 2013 1 2 600 600 0 819
## 9 2013 1 2 600 600 0 846
## 10 2013 1 2 600 600 0 737
## # ... with 2,450 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# between() is a concise alternative for determining if numeric values fall in a range
flights %>% filter(between(dep_time, 600, 605))
```

```
# side note: is.na() can also be useful when filtering
flights %>% filter(!is.na(dep_time))
```

```
# slice() filters rows by position
flights %>% slice(1000:1005)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>        <dbl>   <int>
## 1  2013     1     2     809           810          -1     950
## 2  2013     1     2     810           800          10    1008
## 3  2013     1     2     811           815          -4    1100
## 4  2013     1     2     811           815          -4    1126
## 5  2013     1     2     811           820          -9     944
## 6  2013     1     2     815           815           0    1109
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

```
# keep the first three rows within each group
flights %>%
  group_by(month, day) %>%
  slice(1:3)
```

```
## # A tibble: 1,095 x 19
```

```
## # Groups:   month, day [365]
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
##  1  2013     1     1     517             515         2     830
##  2  2013     1     1     533             529         4     850
##  3  2013     1     1     542             540         2     923
##  4  2013     1     2      42            2359        43     518
##  5  2013     1     2     126            2250       156     233
##  6  2013     1     2     458             500        -2     703
##  7  2013     1     3      32            2359        33     504
##  8  2013     1     3      50            2145       185     203
##  9  2013     1     3     235            2359       156     700
## 10  2013     1     4      25            2359        26     505
## # ... with 1,085 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# sample three rows from each group
flights %>%
  group_by(month, day) %>%
  sample_n(3)
```

```
## # A tibble: 1,095 x 19
## # Groups:   month, day [365]
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
##  1  2013     1     1    1130             1125         5    1301
##  2  2013     1     1    1127             1129        -2    1303
##  3  2013     1     1    1935             1930         5    2223
##  4  2013     1     2    1907             1820        47    2037
##  5  2013     1     2    1648             1635        13    1843
##  6  2013     1     2      NA              1330        NA      NA
##  7  2013     1     3    2005             1925        40    2308
##  8  2013     1     3    2008             1540       268    2339
##  9  2013     1     3    2058             2100        -2    2202
## 10  2013     1     4     909              904         5    1234
## # ... with 1,085 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## # minute <dbl>, time_hour <dtm>
```

```
# keep three rows from each group with the top dep_delay
```

```
flights %>%
```

```
  group_by(month, day) %>%
```

```
  top_n(3, dep_delay)
```

```
## # A tibble: 1,108 x 19
```

```
## # Groups:   month, day [365]
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     1     848           1835         853     1001
```

```
## 2  2013     1     1    1815           1325         290     2120
```

```
## 3  2013     1     1    2343           1724         379      314
```

```
## 4  2013     1     2    1412            838         334     1710
```

```
## 5  2013     1     2    1607           1030         337     2003
```

```
## 6  2013     1     2    2131           1512         379     2340
```

```
## 7  2013     1     3    2008           1540         268     2339
```

```
## 8  2013     1     3    2012           1600         252     2314
```

```
## 9  2013     1     3    2056           1605         291     2239
```

```
## 10 2013     1     4    2058           1730         208        2
```

```
## # ... with 1,098 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## #   minute <dbl>, time_hour <dtm>
```

```
# also sort by dep_delay within each group
```

```
flights %>%
```

```
  group_by(month, day) %>%
```

```
  top_n(3, dep_delay) %>%
```

```
  arrange(desc(dep_delay))
```

```
## # A tibble: 1,108 x 19
```

```
## # Groups:   month, day [365]
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     9     641            900        1301     1242
```

```
## 2  2013     6    15    1432           1935        1137     1607
```

```
## 3  2013     1    10    1121           1635        1126     1239
```

```
## 4  2013     9    20    1139           1845        1014     1457
```

```
## 5 2013 7 22 845 1600 1005 1044
## 6 2013 4 10 1100 1900 960 1342
## 7 2013 3 17 2321 810 911 135
## 8 2013 6 27 959 1900 899 1236
## 9 2013 7 22 2257 759 898 121
## 10 2013 12 5 756 1700 896 1058
## # ... with 1,098 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# unique rows can be identified using unique() from base R
flights %>%
  select(origin, dest) %>%
  unique()
```

```
## # A tibble: 224 x 2
##   origin dest
##   <chr> <chr>
## 1 EWR IAH
## 2 LGA IAH
## 3 JFK MIA
## 4 JFK BQN
## 5 LGA ATL
## 6 EWR ORD
## 7 EWR FLL
## 8 LGA IAD
## 9 JFK MCO
## 10 LGA ORD
## # ... with 214 more rows
```

```
# dplyr provides an alternative that is more "efficient"
flights %>%
  select(origin, dest) %>%
  distinct()
```

```
# side note: when chaining, you don't have to include the parentheses if there are no arguments
flights %>%
  select(origin, dest) %>%
  distinct()
```


4 Adding new variables: mutate, transmute, add_rownames

```
# mutate() creates a new variable (and keeps all existing variables)
flights %>% mutate(speed = distance / air_time * 60)
```

```
## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517             515         2     830
## 2  2013     1     1     533             529         4     850
## 3  2013     1     1     542             540         2     923
## 4  2013     1     1     544             545        -1    1004
## 5  2013     1     1     554             600        -6     812
## 6  2013     1     1     554             558        -4     740
## 7  2013     1     1     555             600        -5     913
## 8  2013     1     1     557             600        -3     709
## 9  2013     1     1     557             600        -3     838
## 10 2013     1     1     558             600        -2     753
## # ... with 336,766 more rows, and 13 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, speed <dbl>
```

```
# transmute() only keeps the new variables
flights %>% transmute(speed = distance / air_time * 60)
```

```
## # A tibble: 336,776 x 1
##   speed
##   <dbl>
## 1  370.
## 2  374.
## 3  408.
## 4  517.
## 5  394.
## 6  288.
## 7  404.
## 8  259.
## 9  405.
## 10 319.
```

```
## # ... with 336,766 more rows
```

```
# example data frame with row names
```

```
mtcars %>% head()
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1   4    4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1   4    4
## Datsun 710     22.8   4  108  93 3.85 2.320 18.61  1  1   4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44  1  0   3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0   3    2
## Valiant        18.1   6  225 105 2.76 3.460 20.22  1  0   3    1
```

```
# add_rownames() turns row names into an explicit variable
```

```
mtcars %>%
```

```
  add_rownames("model") %>%
```

```
  head()
```

```
## Warning: Deprecated, use tibble::rownames_to_column() instead.
```

```
## # A tibble: 6 x 12
```

```
##   model      mpg   cyl  disp    hp  drat    wt  qsec    vs    am  gear  carb
##   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Mazda ~    21     6   160   110   3.9   2.62  16.5     0     1     4     4
## 2 Mazda ~    21     6   160   110   3.9   2.88  17.0     0     1     4     4
## 3 Datsun~   22.8     4   108    93   3.85   2.32  18.6     1     1     4     1
## 4 Hornet~   21.4     6   258   110   3.08   3.22  19.4     1     0     3     1
## 5 Hornet~   18.7     8   360   175   3.15   3.44  17.0     0     0     3     2
## 6 Valiant   18.1     6   225   105   2.76   3.46  20.2     1     0     3     1
```

```
# side note: dplyr no longer prints row names (ever) for local data frames
```

```
mtcars %>% tbl_df()
```

```
## # A tibble: 32 x 11
```

```
##       mpg   cyl  disp    hp  drat    wt  qsec    vs    am  gear  carb
##       <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  21     6   160   110   3.9   2.62  16.5     0     1     4     4
## 2  21     6   160   110   3.9   2.88  17.0     0     1     4     4
## 3 22.8     4   108    93   3.85   2.32  18.6     1     1     4     1
## 4 21.4     6   258   110   3.08   3.22  19.4     1     0     3     1
```

```
## 5 18.7      8 360      175 3.15 3.44 17.0      0      0      3      2
## 6 18.1      6 225      105 2.76 3.46 20.2      1      0      3      1
## 7 14.3      8 360      245 3.21 3.57 15.8      0      0      3      4
## 8 24.4      4 147.      62 3.69 3.19 20        1      0      4      2
## 9 22.8      4 141.      95 3.92 3.15 22.9      1      0      4      2
## 10 19.2     6 168.     123 3.92 3.44 18.3      1      0      4      4
## # ... with 22 more rows
```

5 Grouping and counting: summarise, tally, count, group_size, n_groups, ungroup

```
# summarise() can be used to count the number of rows in each group
flights %>%
  group_by(month) %>%
  summarise(cnt = n())
```

```
## # A tibble: 12 x 2
##   month  cnt
##   <int> <int>
## 1     1 27004
## 2     2 24951
## 3     3 28834
## 4     4 28330
## 5     5 28796
## 6     6 28243
## 7     7 29425
## 8     8 29327
## 9     9 27574
## 10    10 28889
## 11    11 27268
## 12    12 28135
```

```
# tally() and count() can do this more concisely
flights %>%
  group_by(month) %>%
  tally()
flights %>% count(month)
```

```
# you can sort by the count
flights %>%
  group_by(month) %>%
  summarise(cnt = n()) %>%
  arrange(desc(cnt))
```

```
## # A tibble: 12 x 2
##   month    cnt
##   <int> <int>
## 1     7 29425
## 2     8 29327
## 3    10 28889
## 4     3 28834
## 5     5 28796
## 6     4 28330
## 7     6 28243
## 8    12 28135
## 9     9 27574
## 10    11 27268
## 11     1 27004
## 12     2 24951
```

```
# tally() and count() have a sort parameter for this purpose
flights %>%
  group_by(month) %>%
  tally(sort = TRUE)
flights %>% count(month, sort = TRUE)
```

```
# you can sum over a specific variable instead of simply counting rows
flights %>%
  group_by(month) %>%
  summarise(dist = sum(distance))
```

```
## # A tibble: 12 x 2
##   month    dist
##   <int>   <dbl>
## 1     1 27188805
## 2     2 24975509
## 3     3 29179636
```

```
## 4      4 29427294
## 5      5 29974128
## 6      6 29856388
## 7      7 31149199
## 8      8 31149334
## 9      9 28711426
## 10     10 30012086
## 11     11 28639718
## 12     12 29954084
```

```
# tally() and count() have a wt parameter for this purpose
flights %>%
  group_by(month) %>%
  tally(wt = distance)
flights %>% count(month, wt = distance)
```

```
# group_size() returns the counts as a vector
flights %>%
  group_by(month) %>%
  group_size()
```

```
## [1] 27004 24951 28834 28330 28796 28243 29425 29327 27574 28889 27268
## [12] 28135
```

```
# n_groups() simply reports the number of groups
flights %>%
  group_by(month) %>%
  n_groups()
```

```
## [1] 12
```

```
# group by two variables, summarise, arrange (output is possibly confusing)
flights %>%
  group_by(month, day) %>%
  summarise(cnt = n()) %>%
  arrange(desc(cnt)) %>%
  print(n = 40)
```

```
## # A tibble: 365 x 3
```

```

## # Groups:   month [12]
##   month   day   cnt
##   <int> <int> <int>
##  1     11    27 1014
##  2      7    11 1006
##  3      7     8 1004
##  4      7    10 1004
##  5     12     2 1004
##  6      7    18 1003
##  7      7    25 1003
##  8      7    12 1002
##  9      7     9 1001
## 10      7    17 1001
## 11      7    31 1001
## 12      8     7 1001
## 13      8     8 1001
## 14      8    12 1001
## 15      7    22 1000
## 16      7    24 1000
## 17      8     1 1000
## 18      8     5 1000
## 19      8    15 1000
## 20     11    21 1000
## 21      7    15  999
## 22      7    19  999
## 23      7    26  999
## 24      7    29  999
## 25      8     2  999
## 26      8     9  999
## 27     11    22  999
## 28      8    16  998
## 29      7    23  997
## 30      7    30  997
## 31      8    14  997
## 32      7    16  996
## 33      8     6  996
## 34      8    19  996
## 35      9    13  996
## 36      9    26  996

```

```
## 37      9      27    996
## 38      4      15    995
## 39      6      20    995
## 40      6      26    995
## # ... with 325 more rows
```

```
# ungroup() before arranging to arrange across all groups
flights %>%
  group_by(month, day) %>%
  summarise(cnt = n()) %>%
  ungroup() %>%
  arrange(desc(cnt))
```

```
## # A tibble: 365 x 3
##   month   day   cnt
##   <int> <int> <int>
## 1     11    27  1014
## 2      7    11  1006
## 3      7     8  1004
## 4      7    10  1004
## 5     12     2  1004
## 6      7    18  1003
## 7      7    25  1003
## 8      7    12  1002
## 9      7     9  1001
## 10     7    17  1001
## # ... with 355 more rows
```

6 Creating data frames: data_frame

`data_frame()` is a better way than `data.frame()` for creating data frames. Benefits of `data_frame()`:

- You can use previously defined columns to compute new columns.
- It never coerces column types.
- It never munges column names.
- It never adds row names.
- It only recycles length 1 input.
- It returns a local data frame (a `tbl_df`).

```
# data_frame() example
```

```
data_frame(a = 1:6, b = a * 2, c = "string", "d+e" = 1) %>% glimpse()
```

```
## Warning: `data_frame()` is deprecated, use `tibble()`.
```

```
## This warning is displayed once per session.
```

```
## Observations: 6
```

```
## Variables: 4
```

```
## $ a      <int> 1, 2, 3, 4, 5, 6
```

```
## $ b      <dbl> 2, 4, 6, 8, 10, 12
```

```
## $ c      <chr> "string", "string", "string", "string", "string", "string"
```

```
## $ `d+e` <dbl> 1, 1, 1, 1, 1, 1
```

```
# data.frame() example
```

```
data.frame(a = 1:6, c = "string", "d+e" = 1) %>% glimpse()
```

```
## Observations: 6
```

```
## Variables: 3
```

```
## $ a      <int> 1, 2, 3, 4, 5, 6
```

```
## $ c      <fct> string, string, string, string, string, string
```

```
## $ d.e <dbl> 1, 1, 1, 1, 1, 1
```

7 Viewing more output: print, View

```
# specify that you want to see more rows
```

```
flights %>% print(n = 15)
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
```

```
## 1  2013     1     1     517           515           2     830
```

```
## 2  2013     1     1     533           529           4     850
```

```
## 3  2013     1     1     542           540           2     923
```

```
## 4  2013     1     1     544           545          -1    1004
```

```
## 5  2013     1     1     554           600          -6     812
```

```
## 6  2013     1     1     554           558          -4     740
```

```
## 7  2013     1     1     555           600          -5     913
```

```
## 8  2013     1     1     557           600          -3     709
```



```
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## 11 2013 1 1 558 600 -2 849
## 12 2013 1 1 558 600 -2 853
## 13 2013 1 1 558 600 -2 924
## 14 2013 1 1 558 600 -2 923
## 15 2013 1 1 559 600 -1 941
## # ... with 3.368e+05 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# specify that you want to see ALL rows (don't run this!)
flights %>% print(n = Inf)
```

```
# specify that you want to see all columns
flights %>% print(width = Inf)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1 2013     1     1     517           515         2     830
## 2 2013     1     1     533           529         4     850
## 3 2013     1     1     542           540         2     923
## 4 2013     1     1     544           545        -1    1004
## 5 2013     1     1     554           600        -6     812
## 6 2013     1     1     554           558        -4     740
## 7 2013     1     1     555           600        -5     913
## 8 2013     1     1     557           600        -3     709
## 9 2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
##   sched_arr_time arr_delay carrier flight tailnum origin dest  air_time
##             <int>      <dbl> <chr>   <int> <chr>   <chr> <chr>   <dbl>
## 1             819          11 UA      1545 N14228 EWR   IAH     227
## 2             830          20 UA      1714 N24211 LGA   IAH     227
## 3             850          33 AA      1141 N619AA  JFK   MIA     160
## 4            1022         -18 B6       725 N804JB  JFK   BQN     183
## 5             837         -25 DL       461 N668DN  LGA   ATL     116
## 6             728          12 UA      1696 N39463 EWR   ORD     150
## 7             854          19 B6       507 N516JB  EWR   FLL     158
```

```
## 8          723          -14 EV          5708 N829AS LGA      IAD          53
## 9          846          -8 B6           79 N593JB JFK      MCO          140
## 10         745           8 AA           301 N3ALAA LGA      ORD          138
## distance hour minute time_hour
##      <dbl> <dbl>  <dbl> <dtm>
## 1      1400     5      15 2013-01-01 05:00:00
## 2      1416     5      29 2013-01-01 05:00:00
## 3      1089     5      40 2013-01-01 05:00:00
## 4      1576     5      45 2013-01-01 05:00:00
## 5       762     6       0 2013-01-01 06:00:00
## 6       719     5      58 2013-01-01 05:00:00
## 7      1065     6       0 2013-01-01 06:00:00
## 8       229     6       0 2013-01-01 06:00:00
## 9       944     6       0 2013-01-01 06:00:00
## 10      733     6       0 2013-01-01 06:00:00
## # ... with 336,766 more rows
```

```
# show up to 1000 rows and all columns
flights %>% View()

# set option to see all columns and fewer rows
options(dplyr.width = Inf, dplyr.print_min = 6)

# reset options (or just close R)
options(dplyr.width = NULL, dplyr.print_min = 10)
```

8 plot

```
library(ggplot2)

flights %>%
  group_by(dest) %>%
  summarize(
    count = n(),
    dist = mean(distance, na.rm = TRUE),
    delay = mean(arr_delay, na.rm = TRUE)
  ) %>%
  filter(delay, count > 20, dest != "HNL") %>%
```

```
ggplot(mapping = aes(x = dist, y = delay)) +  
  geom_point(aes(size = count), alpha = 1 / 3) +  
  geom_smooth(se = FALSE)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

